Physics task benchmark

10 frozen tasks · High reasoning · 3 attempts per task and model

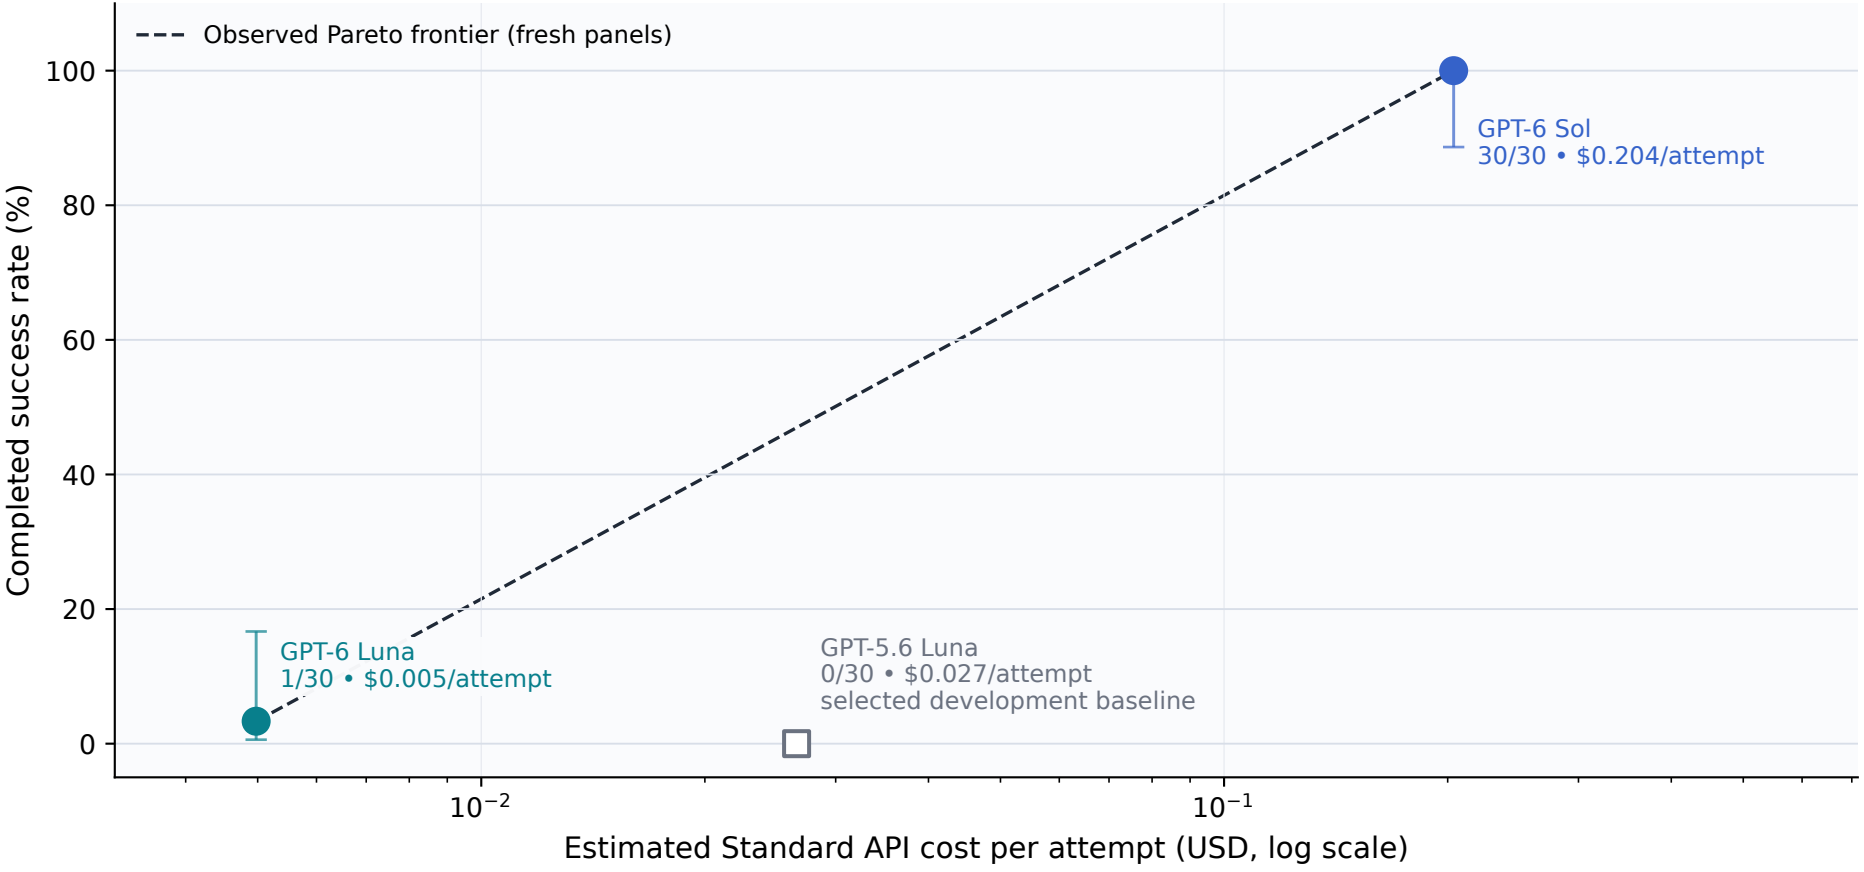

Tokens priced per request, including cache rates; this is not the Pro subscription bill.
Bars: descriptive 95% Wilson intervals (small, task-clustered sample). Tasks were selected against GPT-5.6 Luna.
In progress: GPT-6 Astra. Terra and DeepSeek are not included; model/access details remain unresolved.