Physics benchmark: reasoning effort and cost

13 frozen tasks · 3 attempts per task · 15 complete model–effort points

Stopped at usage guard: DeepSeek None is incomplete (3/39 attempts) and excluded.

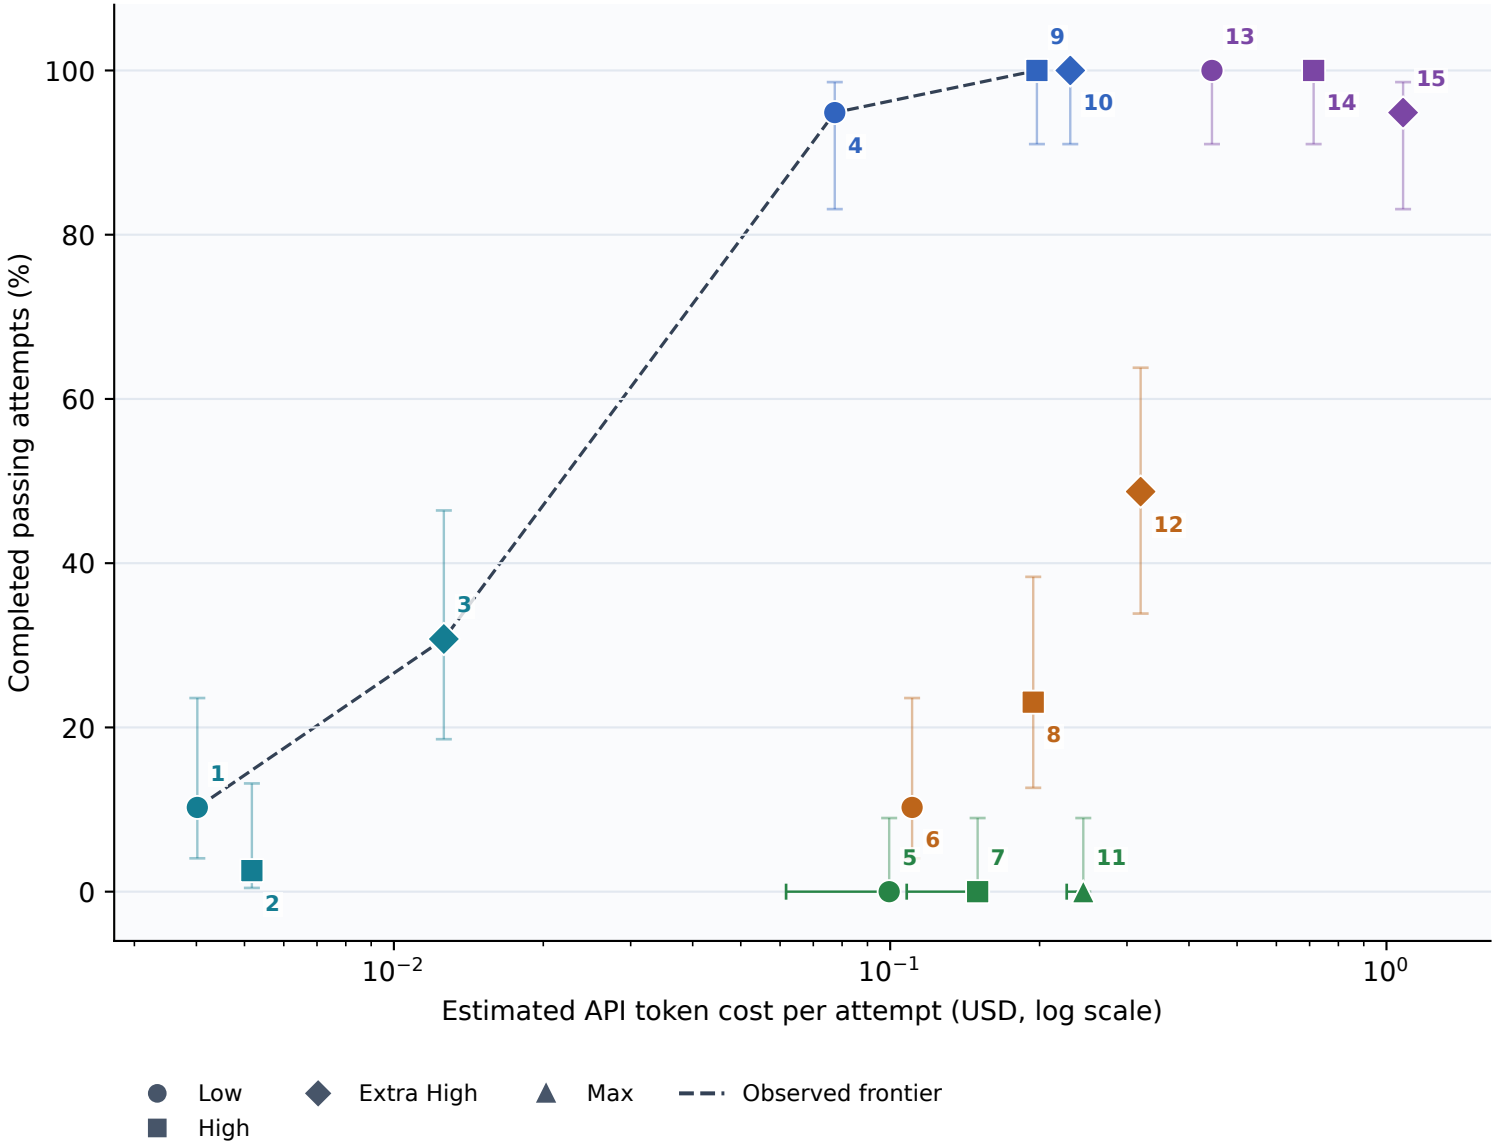

MODEL / EFFORT	SCORE · COST
1 GPT-6 Luna / Low	4/39 · $0.0040
2 GPT-6 Luna / High	1/39 · $0.0052
3 GPT-6 Luna / Extra High	12/39 · $0.0126 · 4 errors
4 GPT-6 Sol / Low	37/39 · $0.0773
5 DeepSeek V4.1 Flash / Low	0/39 · $0.0995 · 33 errors
6 GPT-5.6 Terra / Low	4/39 · $0.1107
7 DeepSeek V4.1 Flash / High	0/39 · $0.1500 · 34 errors
8 GPT-5.6 Terra / High	9/39 · $0.1942
9 GPT-6 Sol / High	39/39 · $0.1975
10 GPT-6 Sol / Extra High	39/39 · $0.2306
11 DeepSeek V4.1 Flash / Max	0/39 · $0.2450 · 38 errors
12 GPT-5.6 Terra / Extra High	19/39 · $0.3196 · 4 errors
13 GPT-6 Astra / Low	39/39 · $0.4450
14 GPT-6 Astra / High	39/39 · $0.7132
15 GPT-6 Astra / Extra High	37/39 · $1.0801 · 2 errors

Per-request token estimates; excludes local compute and subscription billing. Bars: descriptive 95% Wilson intervals. OpenAI High includes reused results; host-sleep effects are documented in the report. Task selection used GPT-5.6 Luna. Incomplete panels excluded. Interrupted responses may omit usage; Prime ranges retain missing-usage reserves.